

CV assignment-03

-S20170010045 D.Nikhil

-S20170010172 V.Hemanth

-S20170010102 N.N.Sri Ram

Image categorization

Question-1

In this task we have to implement image features, train a classifier using the training samples, and then evaluate the classifier on the test set.

We need to implement bag of visual words model and implement Kmeans cluster algorithm to compute visual word dictionary


Steps for implementation:

- 1.Loading labels of training and test data.
- 2.extracting features and regions of given data.
- 3.computing visual bag-of-words dictionary
- 4.implementing k-means clustering algorithm to features.
- 5.calculating confusion matrix and accuracy.

Accuracy is 0.47

Confusion matrix is attached below

precision	recall	f1-score	support		
	1	0.39	0.53	0.45	100
	2	0.69	0.82	0.75	100
	3	0.44	0.31	0.36	100
	4	0.50	0.45	0.47	100
	5	0.37	0.42	0.39	100
	6	0.21	0.27	0.24	100
	7	0.52	0.43	0.47	100
	8	0.44	0.27	0.33	100
accuracy				0.44	800
macro avg		0.45	0.44	0.43	800
weighted avg		0.45	0.44	0.43	800



Scale :

X-Axis : Values of K (Neighbours) from 1 to 100

Y-Axis : Measures of Accuracy

Question-2

1) We are using the pretrained AlexNet. The classifier[6] is changed to 8 to get 8 out_features which was asked to implement.

2) Read the csv files of train and test and create dictionary with image and corresponding labels for it.

3) Initialize the model using initialize_model function where you load the pretrained model of alexnet.

4) We will get the parameters to be optimized or updated. then we will know what parameters to be learned.

5) Training the model and then testing the model and we got 90% accuracy

```
epoch 1 total loss : 381.94107869267464 total correct : 1509 accuracy
: 79.92584745762711
epoch 2 total loss : 323.50325098633766 total correct : 1610 accuracy
: 85.27542372881356
epoch 3 total loss : 247.64012730121613 total correct : 1666 accuracy
: 88.24152542372882
epoch 4 total loss : 265.58933478593826 total correct : 1641 accuracy
: 86.91737288135593
epoch 5 total loss : 284.98790097236633 total correct : 1643 accuracy
: 87.02330508474576
```

epoch 6 total loss : 250.1824003458023 total correct : 1669 accuracy : 88.40042372881356
epoch 7 total loss : 223.3920721411705 total correct : 1693 accuracy : 89.67161016949152
epoch 8 total loss : 187.0485027730465 total correct : 1725 accuracy : 91.36652542372882
epoch 9 total loss : 179.31294015049934 total correct : 1716 accuracy : 90.88983050847457
epoch 10 total loss : 191.97559148073196 total correct : 1712 accuracy : 90.67796610169492
epoch 11 total loss : 135.82300490140915 total correct : 1752 accuracy : 92.79661016949152
epoch 12 total loss : 155.25749826431274 total correct : 1733 accuracy : 91.79025423728814
epoch 13 total loss : 153.37523424625397 total correct : 1741 accuracy : 92.21398305084746
epoch 14 total loss : 165.06358033418655 total correct : 1726 accuracy : 91.41949152542372
epoch 15 total loss : 141.56284427642822 total correct : 1747 accuracy : 92.53177966101694
epoch 16 total loss : 136.85224813222885 total correct : 1747 accuracy : 92.53177966101694
epoch 17 total loss : 129.29159700870514 total correct : 1752 accuracy : 92.79661016949152
epoch 18 total loss : 125.71328336000443 total correct : 1743 accuracy : 92.31991525423729
epoch 19 total loss : 124.69214099645615 total correct : 1764 accuracy : 93.4322033898305
epoch 20 total loss : 132.67240342497826 total correct : 1756 accuracy : 93.00847457627118

time taken to train : 231.4899754524231

```
tensor([[95., 0., 0., 0., 0., 5., 0., 0.],  
        [0., 88., 0., 0., 1., 11., 0., 0.],  
        [0., 0., 96., 0., 0., 2., 2., 0.],  
        [0., 0., 1., 91., 0., 0., 6., 2.],  
        [1., 0., 1., 0., 94., 4., 0., 0.],  
        [4., 1., 1., 0., 2., 92., 0., 0.],  
        [0., 0., 1., 2., 0., 0., 97., 0.],  
        [0., 0., 0., 1., 0., 0., 1., 98.]])
```

Question 3

Steps for implementation:

- 1) At first we have to read images and labels for the data
- 2) Then build the resnet model and train the model and test the data.
- 3) Train accuracy is 93.2
- 4) Test accuracy is 89.32

ResNet-18 Architecture

Layer Name	Output Size	ResNet-18
conv-1	112 x 112 x 64	7 x 7, 64, stride 2
conv-2	56 x 56 x 64	3 x 3 max pool, stride 2 ----- [3 x 3,64] X 2 [3 x 3,64]
conv-3	28 x 28 x 128	[3 x 3,128] X 2 [3 x 3,128]
conv-4	14 x 14 x 256	[3 x 3,256] X 2 [3 x 3,256]
conv-5	7 x 7 x 512	[3 x 3,512] X 2 [3 x 3,512]
Average pool	1 x 1 x 512	7 x 7 average pool
Fully connected	6	512 x 6 fully connections
softmax	6	

Parameters:

Total params: 11,179,590

Trainable params: 11,179,590

Non-trainable params: 0

Total Epochs :70

Loss and Accuracy:

Training Loss: 0.0058141967786823285

Training Accuracy: 94%

Validation Loss: 0.010291938092559577

Validation Accuracy: 89.375%

On cifar10 dataset:

Validation Accuracy -> 77.35%

Comparison:**With KNN**

1)ResNet(DL Model)(with accuracy of around 89%) performs much better than the nearest neighbor classifier(ML Model)(with accuracy of around 49%) in terms of accuracy.

2)ResNet being a deep model was able to learn many complex features and therefore perform better as compared to the knn model, provided we have enough examples to train.

3)ResNet turned out to be the best according to the accuracy.

With Alexnet

Alexnet is a sequential model trained on Imagenet. So it may be the case that not many layers need to be changed as the features learnt for it are found to give high validation accuracy. bggResnet18 is a rather complex model. It seems to overfit into the data given the validation accuracy is much lesser than train accuracy. This particular happens when the model learns complex features unique to the training examples